

The Relational Model

Computer Science Department
Columbia University

Why Study the Relational Model?

- ❖ By far, the most widely used model. Major vendors:
 - Oracle, 44% market share as of end 2011 (source: Gartner)
 - IBM (DB2 and Informix), 22%
 - Microsoft (SQL Server), 18%
 - SAP (bought Sybase in 2010), 4.1%
 - Teradata, 3.3%
- ❖ Other models exist
 - “Legacy systems” in older DBMSs, e.g., IBM’s IMS
 - Object-oriented model (based on some OO language)
 - NoSQL
 - Application specific models (limited use); e.g., event-oriented
 - *Object-Relational model* - most relational DBMS currently employ OO concepts.

Relational Database: Definitions

- ❖ *Relational database*: a set of *relations (tables)*
- ❖ A relation (table) is a set of tuples (rows)
 - *Schema* : specifies the structure: name of relation, name and type of each column, constraints.
 - e.g. Students(*sid* string, *name* string, *login* string, *age* integer, *gpa* real).
 - #columns= *degree / arity*
 - *Instance* : the actual *table* (i.e., its rows) at a particular point in time.
- ❖ All rows in a table are distinct.

Example Instance of Students Relation

sid	name	login	age	gpa
53666	Jones	jones@cs	18	3.4
53688	Smith	smith@eecs	18	3.2
53650	Smith	smith@math	19	3.8

- ❖ All rows distinct
- ❖ Do all columns in a relation instance have to be distinct? (How about all column names?)

Relational Query Languages

- ❖ A major strength of the relational model: supports simple, powerful *querying* of data.
- ❖ Queries can be written intuitively, and the DBMS is responsible for efficient evaluation.
 - The key: precise semantics for relational queries.
 - Allows the optimizer to extensively re-order operations, and still ensure that the answer does not change.

The SQL Query Language

- ❖ Developed by IBM (system R) in the 1970s
- ❖ Need for a standard since it is used by many vendors
- ❖ Standards:
 - SQL-86
 - ...
 - SQL-99 (major extensions including some OO concepts, current standard)

The SQL Query Language

Example:

sid	name	login	age	gpa
53666	Jones	jones@cs	18	3.4
53688	Smith	smith@eecs	18	3.2
53650	Smith	smith@math	19	3.8

❖ To find all 18 year old students, we can write:

```
SELECT *  
FROM Students S  
WHERE S.age=18
```

sid	name	login	age	gpa
53666	Jones	jones@cs	18	3.4
53688	Smith	smith@ee	18	3.2

- To find just names and logins, replace the first line:

```
SELECT S.name, S.login
```

Creating Relations in SQL

- ❖ Creates the Students relation. The type **(domain)** of each field is specified, and enforced by the DBMS whenever tuples are added or modified.

```
CREATE TABLE Students  
(sid CHAR(20),  
name CHAR(20),  
login CHAR(10),  
age INTEGER,  
gpa REAL)
```

- ❖ As another example, the Enrolled table holds information about courses that students take.

```
CREATE TABLE Enrolled  
(sid CHAR(20),  
cid CHAR(20),  
grade CHAR(2))
```


Destroying and Altering Relations

DROP TABLE Students

- ❖ Destroys the relation Students. The schema information *and* the tuples are deleted.

ALTER TABLE Students

ADD COLUMN firstYear: integer

- ❖ The schema of Students is altered by adding a new field; every tuple in the current instance is extended with a *null* value in the new field.

Adding and Deleting Tuples

- ❖ Can insert a single tuple using:

```
INSERT INTO Students (sid, name, login, age, gpa)  
VALUES (53688, 'Smith', 'smith@ee', 18, 3.2)
```

- ❖ Can delete all tuples satisfying some condition (e.g., name = Smith):

```
DELETE  
FROM Students S  
WHERE S.name = 'Smith'
```

** Powerful variants of these commands are available; more later!*

Integrity Constraints (ICs)

- ❖ **IC:** condition that must be true for *any* instance of the database; e.g., domain constraints.
 - ICs are specified when schema is defined.
 - ICs are checked when relations are modified.
- ❖ A *legal* instance of a relation is one that satisfies all specified ICs.
 - DBMS should not allow illegal instances.
- ❖ If the DBMS checks ICs, stored data is more faithful to real-world meaning.
 - Avoids data entry errors, too!

Primary Key Constraints

- ❖ A set of fields is a key for a relation if :
 1. No two distinct tuples can have same values in all key fields, and
 2. This is not true for any subset of the key.
 - Part 2 false? A *superkey*.
 - *Candidate key*: any of the keys
 - *Primary key*: one of the keys chosen DBA
- ❖ e.g., *sid* is a key for Students. *name* is not. The set {*sid*, *gpa*} is a superkey.

SQL – Primary/Unique/Not null

- ❖ **PRIMARY KEY** (att1, att2, ...), only one primary key
UNIQUE (att1, att2, ...)
att **NOT NULL**

- ❖ “For a given student and course, there is a single grade.” **vs.**
“students can take only one course, and receive a single grade for that course; further, no two students in a course receive the same grade.”

```
CREATE TABLE Students (  
    ...  
    age INTEGER NOT NULL)
```

```
CREATE TABLE Enrolled (  
    sid CHAR(20)  
    cid CHAR(20),  
    grade CHAR(2),  
    PRIMARY KEY (sid,cid) )
```

```
CREATE TABLE Enrolled (  
    sid CHAR(20)  
    cid CHAR(20),  
    grade CHAR(2),  
    PRIMARY KEY (sid),  
    UNIQUE (cid, grade) )
```

Foreign Keys, Referential Integrity

- ❖ Foreign key : Set of fields in one tuple that is used to `refer' to another tuple. (Must correspond to primary key of the referenced tuple.) Like a `logical pointer'.
- ❖ Referential integrity is achieved if all foreign key constraints are enforced, i.e., no dangling references.
 - Can you name a data model without referential integrity?
 - Links in HTML!

SQL - Foreign Keys

- ❖ Only students listed in the Students relation should be allowed to enroll for courses.

```
CREATE TABLE Enrolled
(sid CHAR(20), cid CHAR(20), grade CHAR(2),
PRIMARY KEY (sid,cid),
FOREIGN KEY (sid) REFERENCES Students )
```

Enrolled

sid	cid	grade
53666	Carnatic101	C
53666	Reggae203	B
53650	Topology112	A
53666	History105	B

Students

sid	name	login	age	gpa
53666	Jones	jones@cs	18	3.4
53688	Smith	smith@eecs	18	3.2
53650	Smith	smith@math	19	3.8

Enforcing Referential Integrity

- ❖ Consider Students and Enrolled; *sid* in Enrolled is a foreign key that references Students.
- ❖ What should be done if an Enrolled tuple with a non-existent student id is inserted? (*Reject it!*)
- ❖ What should be done if a Students tuple is deleted?
 - Also delete all Enrolled tuples that refer to it.
 - Disallow deletion of a Students tuple that is referred to.
 - Set *sid* in Enrolled tuples that refer to it to a *default sid*.
 - (In SQL, also: Set *sid* in Enrolled tuples that refer to it to a special value *null*, denoting ‘unknown’ or ‘inapplicable’.)
- ❖ Similar if primary key of Students tuple is updated.

SQL - Enforcing referential Integrity

- ❖ SQL/99 supports all 4 options on deletes and updates.
 - Default is **NO ACTION** (*delete/update is rejected*)
 - **CASCADE** (also delete all tuples that refer to deleted tuple)
 - **SET NULL / SET DEFAULT** (sets foreign key value of referencing tuple)

```
CREATE TABLE Enrolled
(sid CHAR(20),
cid CHAR(20),
grade CHAR(2),
PRIMARY KEY (sid,cid),
FOREIGN KEY (sid)
REFERENCES Students
ON DELETE CASCADE
ON UPDATE SET DEFAULT )
```

Where do ICs Come From?

- ❖ ICs are based upon the semantics of the real-world enterprise that is being described in the database relations.
- ❖ We can check a database instance to see if an IC is violated, but we can **NEVER** infer that an IC is true by looking at an instance.
 - An IC is a statement about *all possible* instances!
 - From example, we know *name* is not a key, but the assertion that *sid* is a key is given to us.
- ❖ Key and foreign key ICs are the most common; more general ICs supported too.

Views

- ❖ A view is just a relation, but we store a *definition*, rather than a set of tuples.

```
CREATE VIEW YoungGoodStudents (sid, name, gpa) AS (  
    SELECT S.sid, S.name, E.gpa  
    FROM Students S  
    WHERE S.age < 20 and S.gpa >= 3.0)
```


- ❖ Views can be dropped using the **DROP VIEW** command.
 - How to handle **DROP TABLE** if there's a view on the table?
 - DROP TABLE command has options to let the user specify this.

Views and Security

- ❖ Views can be used to present necessary information (or a summary), while hiding details in underlying relation(s).
 - Given YoungStudents, but not Students or Enrolled, we can find students who are enrolled, but not the *cid*'s of the courses they are enrolled in.


Logical DB Design: ER to Relational

❖ Entity sets to tables:



```
CREATE TABLE Employees  
  (ssn CHAR(11),  
   name CHAR(20),  
   age INTEGER,  
   PRIMARY KEY (ssn))
```

Relationship Sets to Tables




- ❖ In translating a relationship set to a relation, attributes of the relation must include:
 - Keys for each participating entity set (as foreign keys)
 - All descriptive attributes.

```
CREATE TABLE Works_In (  
  ssn CHAR(11),  
  did INTEGER,  
  since DATE,  
  PRIMARY KEY (ssn, did),  
  FOREIGN KEY (ssn)  
    REFERENCES Employees,  
  FOREIGN KEY (did)  
    REFERENCES Departments)
```

Key Constraints

- ❖ Each dept has at most one manager, according to the key constraint on Manages.



*Translation to
relational model?*

Translating ER Diagrams with Key Constraints


1. Map relationship to a table:
 - Note that **did** is the key now!
 - Separate tables for Employees and Departments.
2. Since each department has a unique manager, we could instead combine Manages and Departments.

```
CREATE TABLE Manages (  
  did INTEGER,  
  ssn CHAR(11),  
  since DATE,  
  PRIMARY KEY (did),  
  FOREIGN KEY (ssn) REFERENCES Employees,  
  FOREIGN KEY (did) REFERENCES Departments)
```

```
CREATE TABLE Dept_Mgr (  
  did INTEGER,  
  ssn CHAR(11),  
  since DATE,  
  dname CHAR(20),  
  budget REAL,  
  PRIMARY KEY (did),  
  FOREIGN KEY (ssn) REFERENCES Employees)
```


Participation Constraints

- ❖ Does every department have a manager?
 - If so, this is a participation constraint: the participation of Departments in Manages is said to be *total* (vs. *partial*).
 - Every *did* value in Departments table must appear in a row of the Manages table (with a non-null *ssn* value!)



Participation Constraints in SQL

- ❖ We can capture participation constraints involving one entity set in a binary relationship, but little else (without resorting to CHECK constraints).

```
CREATE TABLE Dept_Mgr(  
  did INTEGER,  
  dname CHAR(20),  
  budget REAL,  
  ssn CHAR(11) NOT NULL,  
  since DATE,  
  PRIMARY KEY (did),  
  FOREIGN KEY (ssn) REFERENCES Employees,  
  ON DELETE NO ACTION)
```


Would this work???

```
CREATE TABLE Manages(  
  ssn CHAR(11) NOT NULL,  
  did INTEGER,  
  since DATE,  
  PRIMARY KEY (did),  
  FOREIGN KEY (ssn) REFERENCES Employees,  
  FOREIGN KEY (did) REFERENCES Departments)
```

Not exactly the same – Why?

Review: Weak Entities

- ❖ A *weak entity* can be identified uniquely only by considering the primary key of another (*owner*) entity.
 - Owner entity set and weak entity set must participate in a one-to-many relationship set (1 owner, many weak entities).
 - Weak entity set must have total participation in this *identifying* relationship set.



Translating Weak Entity Sets

- ❖ Weak entity set and identifying relationship set are translated into a single table.
 - When the owner entity is deleted, all owned weak entities must also be deleted.

```
CREATE TABLE Dep_Policy (  
    pname CHAR(20),  
    age INTEGER,  
    cost REAL,  
    ssn CHAR(11),  
    PRIMARY KEY (pname, ssn),  
    FOREIGN KEY (ssn) REFERENCES Employees  
    ON DELETE CASCADE)
```

A quick note

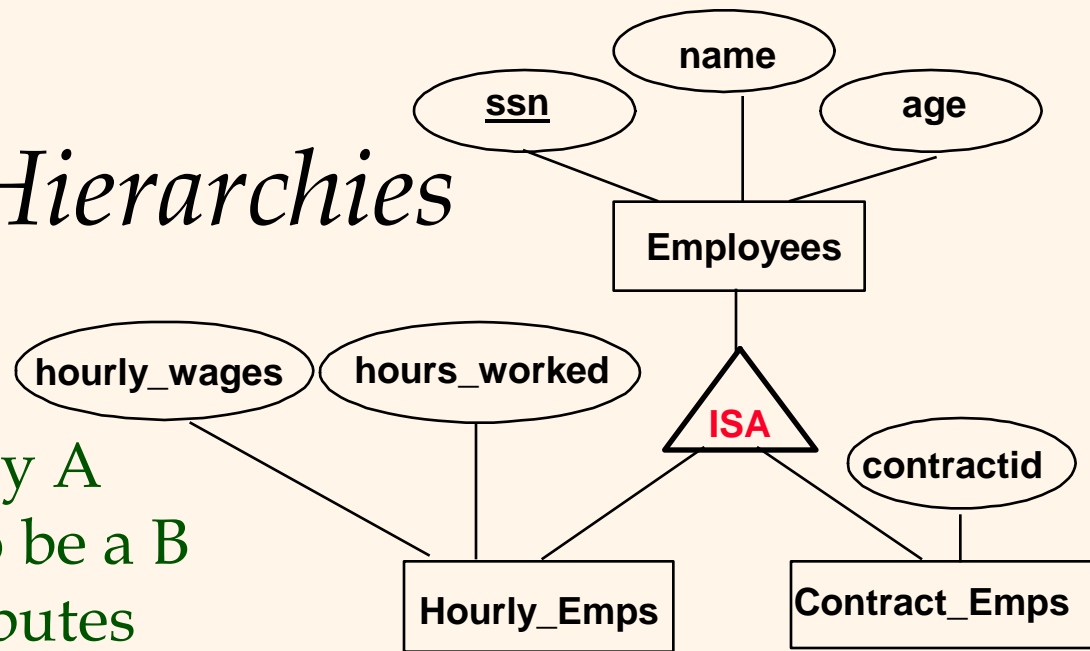
During the initial phases of the DB design process if the attribute types are not that important in clarifying the design, we may write our table definitions in a more compact form by omitting the types; e.g., we may write the previous example as follows:

Dept_Policy (pname, ssn, age, cost, PK(pname, ssn),
FK(ssn) → Employees on delete cascade)

Dept_Policy (pname, ssn, age, cost,
FK(ssn) → Employees on delete cascade)

[But without omitting the **always** crucial PK, UNIQUE, NOT NULL and FK constraints]

Review: ISA Hierarchies



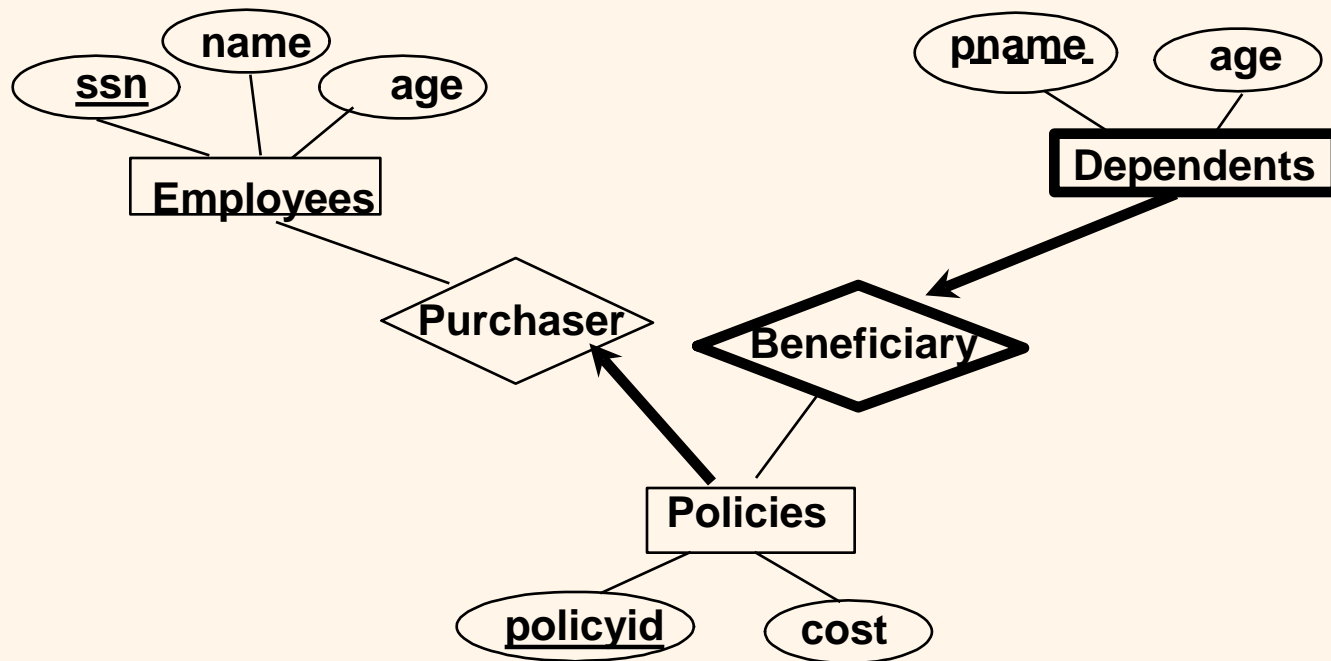
If we declare A **ISA** B, every A entity is also considered to be a B entity. As in OO PLs, attributes are inherited.

- ❖ *Overlap constraints*: Can Joe be an Hourly_Emps as well as a Contract_Emps entity? *(Allowed/disallowed)*
- ❖ *Covering constraints*: Does every Employees entity also have to be an Hourly_Emps or a Contract_Emps entity? *(Yes/no)*

Translating ISA Hierarchies to Relations

1. General approach (one table/set):
 - Emps(ssn, name, age, PK(ssn))
 - Hourly_Emps(h_wages, h_worked, ssn, PK(ssn), FK(ssn) → Emps on delete cascade)
 - Contract_Emps(c_id, ssn, PK(ssn), FK(ssn) → Emps on delete cascade)
 - Queries involving all employees easy; those involving just Hourly_Emps require a join to get some attributes (who does the join?)
 2. Or, may be:
 - Hourly_Emps(ssn, name, age, h_wages, h_worked)
 - Contract_Emps(ssn, name, age, c_id)
 - Queries involving just Hourly_ or Contract_Emps easy; those involving all Emps require a join (by whom?)
- ❖ Effect of overlap/covering constraints?

Multiple relationships - ER diagram



Multiple relationships – SQL mapping

- ❖ The key constraints allow us to combine Purchaser with Policies and Beneficiary with Dependents.


```
CREATE TABLE Policies (  
  policyid INTEGER,  
  cost REAL,  
  ssn CHAR(11) NOT NULL,  
  PRIMARY KEY (policyid).  
  FOREIGN KEY (ssn) REFERENCES Employees,  
  ON DELETE CASCADE)
```

- ❖ Participation constraints lead to NOT NULL constraints.

```
CREATE TABLE Dependents (  
  pname CHAR(20),  
  age INTEGER,  
  policyid INTEGER,  
  PRIMARY KEY (pname, policyid).  
  FOREIGN KEY (policyid) REFERENCES Policies,  
  ON DELETE CASCADE)
```

Orders example - ER diagram

- ❖ Depts, suppliers, parts, orders
- ❖ A department maintains a list of approved suppliers who are the only ones that can supply parts to the department
- ❖ Each supplier maintains a list of parts that can supply



Orders example – SQL mapping

Suppliers (sid, sname)

Parts (pid, pname)

Departments (did, dname)

ApprovedBy (sid, did, FK(sid) → Suppliers, FK(did) → Departments)

Supplies (sid, pid, FK(sid) → Suppliers, FK(pid) → Parts)

Orders (sid, did, pid, qty, cost, date,
FK(sid, did) → ApprovedBy,
FK(sid, pid) → Supplies)

[Choice of PK in Orders has significant implications on what is allowed]

Relational Model: Summary

- ❖ A tabular representation of data.
- ❖ Simple and intuitive, currently the most widely used.
- ❖ Integrity constraints can be specified by the DBA, based on application semantics. DBMS checks for violations.
 - Two important ICs: primary and foreign keys
 - In addition, we *always* have domain constraints.
- ❖ Powerful and natural query languages exist.
- ❖ Rules to translate ER to relational model