

# *SQL: Queries, Programming, Triggers*

Computer Science Department  
Columbia University

## *Reserves*

### *Example Instances*

<u>sid</u>	<u>bid</u>	<u>day</u>
22	101	10/10/96
58	103	11/12/96

## *Sailors*

- ❖ We will use these instances of the Sailors and Reserves relations in our examples.
- ❖ If the key for the Reserves relation contained only the attributes *sid* and *bid*, how would the semantics differ?

<u>sid</u>	sname	rating	age
22	dustin	7	45.0
31	lubber	8	55.5
58	rusty	10	35.0

## *Boats*

<u>bid</u>	bname	color
101	Interlake	blue
102	Interlake	red
103	Clipper	green
104	Marine	red

# Basic SQL Query

SELECT	[DISTINCT] <i>target-list</i>
FROM	<i>relation-list</i>
WHERE	<i>qualification</i>

- ❖ *relation-list* A list of relation names (possibly with a *range-variable* after each name).
- ❖ *target-list* A list of attributes of relations in *relation-list*
- ❖ *qualification* Comparisons (Attr *op* const or Attr1 *op* Attr2, where *op* is one of  $<$ ,  $>$ ,  $=$ ,  $\leq$ ,  $\geq$ ,  $\neq$  ) combined using AND, OR and NOT.
- ❖ **DISTINCT** is an optional keyword indicating that the answer should not contain duplicates. Default is that duplicates are not eliminated!

# Conceptual Evaluation Strategy

- ❖ Semantics of an SQL query defined in terms of the following conceptual evaluation strategy:
  - Compute the cross-product of *relation-list*.
  - Discard resulting tuples if they fail *qualifications*.
  - Delete attributes that are not in *target-list*.
  - If **DISTINCT** is specified, eliminate duplicate rows.
- ❖ Note: the optimization strategy of a particular DBMS may differ but it must come up with *the same answers*.

# Example of Conceptual Evaluation

```
SELECT S.sname  
FROM   Sailors S, Reserves R  
WHERE  S.sid=R.sid AND R.bid=103
```

(sid)	sname	rating	age	(sid)	bid	day
22	dustin	7	45.0	22	101	10/10/96
22	dustin	7	45.0	58	103	11/12/96
31	lubber	8	55.5	22	101	10/10/96
31	lubber	8	55.5	58	103	11/12/96
58	rusty	10	35.0	22	101	10/10/96
58	rusty	10	35.0	58	103	11/12/96

# *A Note on Range Variables*

- ❖ Really needed only if the same relation appears twice in the FROM clause. The previous query can also be written as:

```
SELECT S.sname  
FROM   Sailors S, Reserves R  
WHERE  S.sid=R.sid AND bid=103
```

OR

```
SELECT sname  
FROM   Sailors, Reserves  
WHERE  Sailors.sid=Reserves.sid  
       AND bid=103
```

*It is good style,  
however, to use  
range variables  
most of the time!*

*Find sailors who've reserved at least one boat*

```
SELECT S.sid  
FROM Sailors S, Reserves R  
WHERE S.sid=R.sid
```

- ❖ Would adding DISTINCT to this query make a difference?
- ❖ What is the effect of replacing *S.sid* by *S.sname* in the SELECT clause? Would adding DISTINCT to this variant of the query make a difference?

# *Expressions and Strings*

```
SELECT S.age, age1=S.age-5, 2*S.age AS age2  
FROM Sailors S  
WHERE S.sname LIKE 'B_%B'
```

- ❖ Illustrates use of arithmetic expressions and string pattern matching
- ❖ **AS** and **=** are two ways to name fields in result.
- ❖ **LIKE** is used for string matching. **`\_`** stands for any one character and **`%`** stands for 0 or more arbitrary characters.



*Find sid's of sailors who've reserved a red or a green boat*

- ❖ **UNION**: Can be used to compute the union of any two *union-compatible* sets of tuples (which are themselves the result of SQL queries).
- ❖ **EXCEPT/MINUS**: set difference.
  - What do we get if we replace **UNION** by **EXCEPT**?

```
SELECT R.sid
FROM Boats B, Reserves R
WHERE R.bid=B.bid
      AND (B.color='red' OR B.color='green')
```

*Q: what do we get if we replace OR by AND?*


```
SELECT R.sid
FROM Boats B, Reserves R
WHERE R.bid=B.bid
      AND B.color='red'
UNION
SELECT R.sid
FROM Boats B, Reserves R
WHERE R.bid=B.bid
      AND B.color='green'
```

*Find sid's of sailors who've reserved a red and a green boat*

```
SELECT R1.sid
FROM Boats B1, Reserves R1,
     Boats B2, Reserves R2
WHERE R1.sid=R2.sid AND R1.bid=B1.bid
                        AND R2.bid=B2.bid
AND (B1.color='red' AND B2.color='green')
```

- ❖ **INTERSECT**: Can be used to compute the intersection of any two *union-compatible* sets of tuples. (Included in the SQL/92 standard, but some systems don't support it.)
- ❖ Contrast symmetry of the UNION and INTERSECT queries with how much the other versions differ.

```
SELECT R.sid
FROM Boats B, Reserves R
WHERE R.bid=B.bid
      AND B.color='red'
INTERSECT
SELECT R.sid
FROM Boats B, Reserves R
WHERE R.bid=B.bid
      AND B.color='green'
```

 Key field!

# Nested Queries

*Find names of sailors who've reserved boat #103:*

```
SELECT S.sname
```

```
FROM Sailors S
```

```
WHERE S.sid IN (SELECT R.sid  
                FROM Reserves R  
                WHERE R.bid=103)
```


- ❖ A very powerful feature of SQL: a WHERE clause can itself contain an SQL query! (Actually, so can FROM and HAVING clauses; more later.)
- ❖ To find sailors who've *not* reserved #103, use **NOT IN**.
- ❖ To understand semantics of nested queries, think of a nested loops evaluation: *For each Sailors tuple, check the qualification by computing the subquery.*

# *Example in class*

## *Write SQL using IN*

*Find names of sailors who've reserved a red boat:*

```
SELECT S.sname
FROM Sailors S
WHERE S.sid IN (SELECT R.sid
                FROM Reserves R
                WHERE R.bid IN (SELECT B.bid
                               FROM Boats B
                               WHERE B.color=red))
```



How about if we change the query to be:

*Find names of sailors who've not reserved a red boat*

**NOT IN**

# Nested Queries with Correlation

*Find names of sailors who've reserved boat #103:*

```
SELECT S.sname
FROM Sailors S
WHERE EXISTS (SELECT *
              FROM Reserves R
              WHERE R.bid=103 AND S.sid=R.sid)
```



- ❖ **EXISTS** is another set comparison operator, like **IN**; it checks for empty set: returns true if table is not empty
  - The example illustrates why, in general, subquery must be re-computed for each Sailors tuple.
- ❖ **UNIQUE** checks for duplicate tuples: returns true if there are no duplicate tuples or it is empty)
  - If **UNIQUE** is used (instead of EXISTS), and \* is replaced by *R.bid*, finds sailors with at most one reservation for boat #103. (\* denotes all attributes; Why do we have to replace \* by *R.bid*?)

# More on Set-Comparison Operators

- ❖ We've already seen IN, EXISTS and UNIQUE. Can also use NOT IN, NOT EXISTS and NOT UNIQUE.
- ❖ Also available: *op* ANY, *op* ALL, *op* one of >, <, =, ≥, ≤, ≠
- ❖ Find sailors whose rating is greater than that of some sailor called Mike:

```
SELECT *  
FROM Sailors S  
WHERE S.rating > ANY (SELECT S2.rating  
                      FROM Sailors S2  
                      WHERE S2.sname='Mike')
```

# Rewriting INTERSECT Queries Using IN

*Find id/name of sailors who've reserved both a red and a green boat:*

```
SELECT S.sid, S.sname
FROM Sailors S, Boats B, Reserves R
WHERE S.sid=R.sid AND R.bid=B.bid AND B.color='red'
      AND S.sid IN (SELECT S2.sid
                    FROM Sailors S2, Boats B2, Reserves R2
                    WHERE S2.sid=R2.sid AND R2.bid=B2.bid
                      AND B2.color='green')
```

❖ Similarly, EXCEPT queries re-written using NOT IN.

# Division in SQL

Find sailors who've reserved all boats.

❖ Let's also do it the hard way, without EXCEPT:

(2) SELECT S.sname  
FROM Sailors S

WHERE NOT EXISTS (SELECT B.bid  
FROM Boats B

*Sailors S such that ...* WHERE NOT EXISTS (SELECT R.bid  
FROM Reserves R  
*there is no boat B without ...* WHERE R.bid=B.bid  
AND R.sid=S.sid))  
*a Reserves tuple showing S reserved B*

(1)

```
SELECT S.sname
FROM Sailors S
WHERE NOT EXISTS
    ((SELECT B.bid
      FROM Boats B)
  EXCEPT
  (SELECT R.bid
   FROM Reserves R
  WHERE R.sid=S.sid))
```



# Aggregate Operators

- ❖ Significant extension of relational algebra.

COUNT (\*)  
COUNT ( [DISTINCT] A)  
SUM ( [DISTINCT] A)  
AVG ( [DISTINCT] A)  
MAX (A)  
MIN (A)

*single column*

```
SELECT COUNT (*)  
FROM Sailors S
```

```
SELECT AVG (S.age)  
FROM Sailors S  
WHERE S.rating=10
```

```
SELECT AVG ( DISTINCT S.age)  
FROM Sailors S  
WHERE S.rating=10
```

```
SELECT S.sname  
FROM Sailors S  
WHERE S.rating= (SELECT MAX(S2.rating)  
FROM Sailors S2)
```

## *Find name and age of the oldest sailor(s)*

- ❖ The first query is illegal! (We'll look into the reason a bit later, when we discuss **GROUP BY**.)
- ❖ The third query is equivalent to the second query, and is allowed in the SQL/92 standard, but is not supported in some systems.

```
SELECT S.sname, MAX (S.age)
FROM Sailors S
```

```
SELECT S.sname, S.age
FROM Sailors S
WHERE S.age =
      (SELECT MAX (S2.age)
       FROM Sailors S2)
```

```
SELECT S.sname, S.age
FROM Sailors S
WHERE (SELECT MAX (S2.age)
       FROM Sailors S2)
      = S.age
```

# *GROUP BY and HAVING*

- ❖ So far, we've applied aggregate operators to all (qualifying) tuples. Sometimes, we want to apply them to each of several *groups* of tuples.
- ❖ Consider: *Find the age of the youngest sailor for each rating level.*
  - In general, we don't know how many rating levels exist, and what the rating values for these levels are!
  - Suppose we know that rating values go from 1 to 10; we can write 10 queries that look like this (!):

For  $i = 1, 2, \dots, 10$ :

```
SELECT MIN (S.age)
FROM Sailors S
WHERE S.rating = i
```

# Queries With GROUP BY and HAVING

SELECT	[DISTINCT] <i>target-list</i>
FROM	<i>relation-list</i>
WHERE	<i>qualification</i>
GROUP BY	<i>grouping-list</i>
HAVING	<i>group-qualification</i>

- ❖ The *target-list* contains (i) attribute names (ii) terms with aggregate operations (e.g., MIN (*S.age*)).
  - The attribute list (i) must be a subset of *grouping-list*.  
Intuitively, each answer tuple corresponds to a *group*, and these attributes must have a single value per group. (A *group* is a set of tuples that have the same value for all attributes in *grouping-list*.)

# Conceptual Evaluation

- ❖ The cross-product of *relation-list* is computed, tuples that fail *qualification* are discarded, 'unnecessary' fields are deleted, and the remaining tuples are partitioned into groups by the value of attributes in *grouping-list*.
- ❖ The *group-qualification* is then applied to eliminate some groups. Expressions in *group-qualification* must have a *single value per group!*
  - In effect, an attribute in *group-qualification* that is not an argument of an aggregate op also appears in *grouping-list*. (SQL does not exploit primary key semantics here!)
- ❖ One answer tuple is generated per qualifying group.

*Find the age of the youngest sailor with age > 18, for each rating with at least 2 such sailors*

```
SELECT S.rating, MIN (S.age)
FROM Sailors S
WHERE S.age > 18
GROUP BY S.rating
HAVING COUNT (*) > 1
```

<u>sid</u>	sname	rating	age
22	dustin	7	45.0
31	lubber	8	55.5
71	zorba	10	16.0
64	horatio	7	35.0
29	brutus	1	33.0
58	rusty	10	35.0

rating	age
1	33.0
7	45.0
7	35.0
8	55.5
10	35.0

rating	
7	35.0

*Answer relation*

- ❖ Only S.rating and S.age are mentioned in the SELECT, GROUP BY or HAVING clauses; other attributes *'unnecessary'*.
- ❖ 2nd column of result is unnamed. (Use AS to name it.)

*Find the age of the youngest sailor with age > 18,  
for each rating with at least 2 sailors (of any age)*

```
SELECT S.rating, MIN (S.age)
FROM Sailors S
WHERE S.age > 18
GROUP BY S.rating
HAVING 1 < (SELECT COUNT (*)
             FROM Sailors S2
             WHERE S.rating=S2.rating)
```

- ❖ Shows HAVING clause can also contain a subquery.
- ❖ Compare this with the query where we considered only ratings with 2 sailors over 18!
- ❖ What if HAVING clause is replaced by:
  - HAVING COUNT(\*) >1

*For each red boat, find the number of reservations for this boat*

```
SELECT B.bid, COUNT (*) AS num_of_reservations
FROM Boats B, Reserves R
WHERE R.bid=B.bid AND B.color='red'
GROUP BY B.bid
```

- ❖ Grouping over a join of two relations.
- ❖ What do we get if we remove *B.color='red'* from the WHERE clause and add a HAVING clause with this condition?

Only columns that appear in GROUP BY can appear in HAVING, unless they appear as arguments to an aggregate operator in the HAVING clause



*Find those ratings for which the average age is the minimum over all ratings*

❖ Aggregate operations cannot be nested! **WRONG:**

```
SELECT S.rating
FROM Sailors S
WHERE S.age = (SELECT MIN (AVG (S2.age))
               FROM Sailors S2 GROUP BY S2.rating)
```

❖ Correct solution:

```
SELECT Temp.rating, Temp.avg_age
FROM (SELECT S.rating, AVG (S.age) AS avg_age
      FROM Sailors S
      GROUP BY S.rating) Temp
WHERE Temp.avg_age =
      (SELECT MIN (Temp.avg_age) FROM Temp)
```

# Null Values

- ❖ Field values in a tuple are sometimes *unknown* (e.g., a rating has not been assigned) or *inapplicable* (e.g., no spouse's name).
  - SQL provides a special value *null* for such situations.
- ❖ The presence of *null* complicates many issues; e.g.:
  - Special operators needed to check if value is/is not *null*.
  - Is *rating* > 8 true or false when *rating* is equal to *null*? What about *AND*, *OR* and *NOT* connectives?
  - We need a 3-valued logic (true, false and *unknown*).
  - Meaning of constructs must be defined carefully. (e.g., WHERE clause eliminates rows that don't evaluate to true.)
  - New operators (in particular, *outer joins*) possible/needed.

# Outer Joins

- ❖ Print the sid of every sailor and the boats the sailor has reserved. Join will not work; sailors with no reservations will not be included.

```
SELECT S.sid, R.bid
FROM Sailors S, Reserves R
WHERE S.sid=R.sid
```

sid	bid
22	101
58	103

- ❖ Left Outer Join: rows in the left table that do not match some row in the right table appear exactly once, with column from the right table assigned NULL values.

```
SELECT S.sid, R.bid
FROM Sailors S LEFT OUTER JOIN Reserves R ON S.sid = R.sid
```

sid	bid
22	101
31	<i>null</i>
58	103

- ❖ Right Outer Join: the reverse
- ❖ Full Outer Join: both, left and right
- ❖ Another example: Print the sid of every sailor and the number of boats the sailor has reserved.

```
SELECT S.sid, COUNT(R.bid)
FROM Sailors S LEFT OUTER JOIN Reserves R ON S.sid = R.sid
GROUP BY S.sid
```

# *Integrity Constraints (Review)*

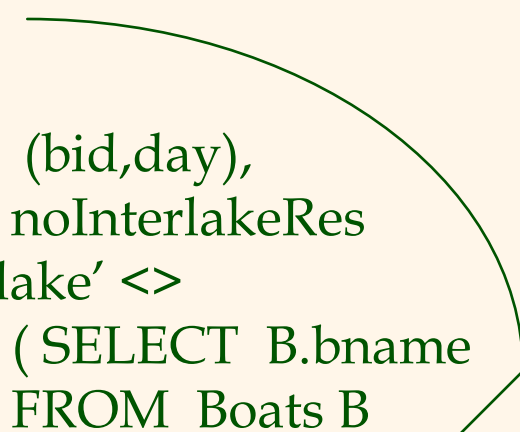
- ❖ Integrity constraints describe conditions that every *legal instance* of a relation must satisfy. It's the DBMS that enforces them on every update – those updates that violates them are disallowed. Types we have seen so far:
  - Domain constraints (field type)
  - Primary key
  - Unique and not null
  - Foreign keys
- ❖ We will now discuss general constraints.

# General Constraints (attribute- & tuple-based)

- ❖ Useful when more general ICs than keys are involved.
- ❖ Can use queries to express constraint.
- ❖ Constraints can be named.
- ❖ Required to hold on nonempty tables only.

```
CREATE TABLE Sailors
( sid INTEGER,
  sname CHAR(10),
  rating INTEGER,
  age REAL,
  PRIMARY KEY (sid),
  CHECK ( rating >= 1
        AND rating <= 10 )
```

```
CREATE TABLE Reserves
( sname CHAR(10),
  bid INTEGER,
  day DATE,
  PRIMARY KEY (bid,day),
  CONSTRAINT noInterlakeRes
  CHECK (`Interlake' <>
        ( SELECT B.bname
          FROM Boats B
          WHERE B.bid=bid)))
```



# Constraints Over Multiple Relations

```
CREATE TABLE Sailors
```

```
( sid INTEGER,  
  sname CHAR(10),  
  rating INTEGER,  
  age REAL,  
  PRIMARY KEY (sid),  
  CHECK
```

*Number of boats  
plus number of  
sailors is < 100*

- ❖ Awkward and wrong!
- ❖ If Sailors is empty, the number of Boats tuples can be anything!
- ❖ **ASSERTION** is the right solution; not associated with either table.

```
( (SELECT COUNT (S.sid) FROM Sailors S)  
  + (SELECT COUNT (B.bid) FROM Boats B) < 100 )
```

```
CREATE ASSERTION smallClub
```

```
CHECK
```

```
( (SELECT COUNT (S.sid) FROM Sailors S)  
  + (SELECT COUNT (B.bid) FROM Boats B) < 100 )
```

# *Using assertions to enforce total participation*

- ❖ All sailors must reserve at least one boat (i.e., total participation of Sailors in Reserves)

```
CREATE ASSERTION BusySailors
CHECK (
  NOT EXISTS (
    SELECT sid FROM Sailors
    WHERE sid NOT IN (
      SELECT sid FROM Reserves)
    )
)
```

# *Triggers*

- ❖ Trigger: procedure that starts automatically if specified changes occur to the DBMS
- ❖ Three parts:
  - Event (activates the trigger)
  - Condition (tests whether the triggers should run)
  - Action (what happens if the trigger runs)



# *Triggers Issues*

- ❖ Before (a change occurs) vs after
- ❖ When a statement triggers an event, execution of the trigger's action may occur in one of the following fashion
  - Instead
  - Deferred
  - Asynchronous
- ❖ Transactions
  - Part of the statement vs independent action

# *Triggers Example*

```
CREATE TRIGGER init_count BEFORE INSERT ON Students
  DECLARE count INTEGER;
  BEGIN
    count := 0;
  END
```

```
CREATE TRIGGER incr_count AFTER INSERT ON Students
  WHEN (new.age < 18) /* also, old (before value) */
  FOR EACH ROW
  BEGIN
    count := count + 1;
  END
```

# *Summary*

- ❖ SQL was an important factor in the early acceptance of the relational model; more natural than earlier, procedural query languages.
- ❖ Relationally complete; in fact, significantly more expressive power than relational algebra.
- ❖ Even queries that can be expressed in RA can often be expressed more naturally in SQL.
- ❖ Many alternative ways to write a query; optimizer should look for most efficient evaluation plan.
  - In practice, users need to be aware of how queries are optimized and evaluated for best results.

## *Summary (Contd.)*

- ❖ NULL for unknown field values brings many complications
- ❖ SQL allows specification of rich integrity constraints
- ❖ Triggers respond to changes in the database